

# **Introduction to Modern Databases**

Redwan Ahmed Rizvee

(<https://rizveeredwan.github.io>)

# Modern Database Systems: Concepts & Extensions

- An overview of advanced database concepts including
  - object-oriented,
  - object-relational, and
  - specialized databases such as : *temporal, spatial, multimedia, and mobile databases.*
  - Complex data types: structured, array, multiset
  - Inheritance, object identity, and reference types
  - Object-relational queries and implementation
  - Persistent programming languages

# Introduction to Modern Databases

- Traditional relational databases **are no longer sufficient** for many modern applications.
- The rise of OOP, multimedia, mobile devices, and time/location-sensitive data demands more advanced models.
- Modern DBs extend relational models with richer types, object orientation, persistence, and support for temporal, spatial, multimedia data.

# Object-Oriented Databases (OODB)

- Built on object-oriented programming principles.
- **Key features:**
  - Encapsulation: Data + methods encapsulated within objects.
  - Inheritance: Subtypes inherit attributes and methods from supertypes.
  - Object identity: Every object has a unique, system-generated identifier independent of its values.
- **Advantages:**
  - Closer alignment with application programming models.
  - Improved support for complex and user-defined data types.
  - Reusability and extensibility through class hierarchies.

# OODB Example Scenario

- **Domain:** University Course Management System
- **Classes:**
  - Person { ID, name, address }
  - Student extends Person { major, GPA }
  - Instructor extends Person { department, salary }
  - Course { courseID, title, instructor (ref), enrolledStudents (list of refs) }
- **Relationships:**
  - Inheritance: Student and Instructor inherit from Person.
  - Associations via references (e.g., instructor assigned to a course).

# Query Example in OODB

- **Goal:** List the names of students enrolled in "Database Systems".
- **OQL (Object Query Language):**

```
SELECT s.name  
FROM Course c, c.enrolledStudents s  
WHERE c.title = "Database Systems";
```

- **Interpretation:** Navigates from course to associated students using **path expressions** and **object references**.

```
for course in Course:  
    for student in  
course.enrolledStudents:  
        if course.title == "Database  
Systems":  
            print(student.name)
```

This is not cross!

Course **c**: Iterates over all Course objects.

**c.enrolledStudents s**: For each course **c**, it iterates over the list/collection **enrolledStudents** and binds each student object to **s**.

# Implementation Aspects of OODB

**Storage:** Objects are stored with OIDs and can refer to each other directly.

**Indexing:** Path indexes support efficient access to nested or referenced attributes.

**Persistence:** Objects remain in database across program executions without needing serialization.

# Indexing in OODB

## ***Indexing: Path Indexes Support Efficient Access to Nested or Referenced Attributes***

- In object-oriented databases (OODB), **data is stored as objects that can reference other objects**, often in nested structures or via pointers.
- When querying, you may want **to quickly find objects based on attributes that are not directly in the root object, but inside a referenced object or nested deep within the structure.**
- **Path indexing creates indexes not just on flat attributes, but on paths through these object references.**

# How Indexing Helps?

- a) Suppose you have an object **Order** that references a **Customer** object, which has an attribute **city**.
- b) A path index **could index the path** **Order.customer.city**.
- c) When **you query orders where the customer lives in "New York,"** the database *uses the path index to quickly find all matching orders without scanning all orders and then fetching the customer info.*

# Persistence in OODB

***Persistence: Objects Remain in Database Across Program Executions Without Needing Serialization***

- Persistence means that data (objects) outlive the program that created or modified them — they are stored permanently until explicitly deleted.
- In traditional programming, to save objects to disk or send them over a network, you typically serialize them (convert them into a flat format like JSON or binary), then later deserialize back to objects.
- **In OODBs, persistence is transparent.**

# Tabular Data vs OODB

Category	Tabular Data	OODB
<b>Data Representation</b>	Data is stored in flat tables with rows and columns.	Data is stored as objects with attributes and methods.
<b>Relationships</b>	Represented using foreign keys.	Represented using direct object references.
<b>Complexity Handling</b>	Requires joins for complex relationships.	Uses navigational access and encapsulation.
<b>Schema Evolution</b>	Altering schemas may require data migration.	Easier to evolve due to inheritance and class structure.
<b>Use Case Fit</b>	Well-suited for structured, transactional data.	Ideal for complex data and close integration with object-oriented applications.

# Object-Relational Databases (ORDB)

- Merge of **relational (tabular)** and object-oriented (OODB) models.
- Supported by SQL:1999, SQL:2003.
- Features:
  - User-defined types (UDTs)
  - Table inheritance (subtypes)
  - References (Object Pointers)
  - Structured and collection types
  - Methods on types
- Examples: PostgreSQL, Oracle, IBM DB2

Object-Relational Queries allow you to manipulate and retrieve data **using these extended features while maintaining compatibility with SQL**.

# Example Schema of ORDB

```
-- Define a user-defined type for Address  
CREATE TYPE address_type AS (  
    street VARCHAR,  
    city VARCHAR,  
    zip VARCHAR  
) ;
```

```
-- Define a table Instructor that inherits from Person
```

```
CREATE TABLE Instructor (  
    department VARCHAR,  
    salary NUMERIC  
) INHERITS (Person) ;
```

```
-- Define a base table for Person
```

```
CREATE TABLE Person (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR,  
    addr address_type  
) ;
```

```
-- Table Course with a reference to Instructor
```

```
CREATE TABLE Course (  
    course_id SERIAL PRIMARY KEY,  
    title VARCHAR,  
    instructor_id INT REFERENCES Instructor(id)  
) ;
```

```
-- Define a table Student that inherits from Person
```

```
CREATE TABLE Student (  
    major VARCHAR,  
    gpa FLOAT  
) INHERITS (Person) ;
```

Mainly Postgres Supports this, not all.

# Example Query of ORDB

Relational approach

```
SELECT s.name, s.major, c.title  
FROM Student s  
JOIN Enrollment e ON s.id = e.student_id  
JOIN Course c ON e.course_id = c.course_id  
JOIN Instructor i ON c.instructor_id = i.id  
WHERE s.major = 'Computer Science' AND i.department = 'CS';
```

- we combine inheritance (Person), UDTs (Address) and References (instructor)

```
SELECT DEREF(s).name, c.title  
FROM Course c, UNNEST(c.enrolledStudents) AS s  
WHERE c.title = 'Database Systems';
```

Object-Relational approach

# Key Features of Object Relational Queries

- **Path Expressions:**

Navigate through nested attributes or referenced objects using dot notation.

E.g., `student.address.city`

- **Dereferencing References:**

Objects can have attributes that are references (`REF`) to other objects. You can dereference these in queries.

E.g., `student.advisor.name` — `advisor` is a reference to an `Instructor` object.

For OODBMS, `student.advisor.name` However, for ORDBMS, need to DREF  
`DEREF(student.advisor).name`

- **Table Inheritance Queries:**

Queries can exploit inheritance hierarchies in tables.

E.g., Querying a base table returns all rows including those in child tables.

SQL: `SELECT * FROM Person` returns rows from `Person`, `Student`, and `Instructor` tables if they inherit from `Person`.

- **User-Defined Types and Methods:**

You can call methods (functions) on objects in queries.

E.g., `student.getAge()` to compute the student's age dynamically.

- **Polymorphic Queries:**

Queries can return results of various subtypes, with the ability to discriminate based on actual type.

# Common Syntax Extensions in ORDB Queries

Feature	Example	Description
Path Expression	<code>student.address.city</code>	Access nested fields
Dereference Reference	<code>student.advisor.name</code> <code>DREF(student.advisor).name</code>	Follow reference links
Table Inheritance	<code>SELECT * FROM Person</code>	Retrieve rows from all subtypes
Polymorphic Filter	<code>WHERE typeOf(person) = 'Student'</code>	Filter by subtype
Method Call	<code>student.getGPA()</code>	Invoke methods on objects
Collection Queries	<code>SELECT s FROM course.enrolledStudents s</code>	Query elements of collection attributes

# New Primitives for Object-Relational Queries

1	User-Defined Types (UDTs)
	<pre>CREATE TYPE address_type AS (     street VARCHAR,     city VARCHAR,     zip VARCHAR );</pre>
2	References ( <a href="#">REF</a> ) and Object Identity ( <a href="#">OID</a> ) ( <a href="#">advisor</a> is a <a href="#">REF</a> to an Instructor object)
	<pre>SELECT s.name, DREF(s.advisor).name FROM Student s;</pre>
3	Path Expressions Navigate nested attributes or referenced objects
	<pre>SELECT DEREF(c.instructor).department FROM Course c;</pre>
4	Inheritance and Type Hierarchies  Tables can inherit from base tables Polymorphic queries on super- and subtypes
	<pre>SELECT * FROM Person; -- Returns Persons, Students, Instructors</pre>

# New Primitives for Object-Relational Queries

5	<p>Collection Types Attributes can be arrays, lists, multisets Query elements in collections</p>
	<pre>SELECT DEREF(s).name FROM Course c, UNNEST(c.enrolledStudents) AS s; // UNNEST() is used to flatten collection-type attributes (e.g., arrays, multisets, or lists) into individual rows, allowing relational-style querying over collections stored within tuples.</pre>
6	<p>Methods / Member Functions Functions bound to types, invoked in queries</p>
	<pre>SELECT s.name, s.getGPA() FROM Student s WHERE s.getGPA() &gt; 3.5;</pre>
7	<p>Polymorphic Queries and Type Checking</p>
	<pre>SELECT p.name FROM Person p WHERE TYPEOF(p) = 'Student';</pre>
8	<p>Nested and Correlated Queries</p>
	<pre>SELECT c.title FROM Course c WHERE EXISTS (   SELECT 1 FROM UNNEST(c.enrolledStudents) s WHERE s.gpa &gt; 3.8 );</pre>

# New Primitives for Object-Relational Queries

9	Extended Join Operations
	<pre>SELECT s.name, c.title FROM Student s JOIN Enrollment e ON s.id = e.student_id JOIN Course c ON e.course_id = c.course_id;</pre>
10	Object Creation and Manipulation
	<pre>INSERT INTO Student VALUES (NEW Student(...));</pre>

# Complex Data Types

```
CREATE TYPE Address AS (
    street VARCHAR(50),
    city VARCHAR(30),
    zip VARCHAR(10)
);
```

```
CREATE TABLE Employee (
    empID INT,
    name VARCHAR(50),
    address Address
);
```

Structured Types: Nested records/tuples.  
Used when you want to encapsulate multiple related fields as a single unit.

```
CREATE TABLE Product (
    productID INT,
    name VARCHAR(50),
    tags TEXT[] -- array of tags
);
```

Arrays: Indexed collections of fixed or variable length.  
Useful in scientific data where you store series of measurements.

```
CREATE TYPE IntBag AS MULTISET(INTEGER);
```

Multisets: Unordered collections allowing duplicates.

# Inheritance and Object Identity

Inheritance: Enables polymorphism in queries.

These denotes fields

```
CREATE TYPE Person AS (name TEXT) NOT FINAL;  
CREATE TYPE Student UNDER Person AS (major TEXT);  
CREATE TYPE Instructor UNDER Person AS (department TEXT);
```

```
CREATE TABLE People OF Person; // This creates a table named People where each row is an object of type Person or any of its subtypes (Student or Instructor).
```

## Insertion Query

```
INSERT INTO People VALUES (Student('Alice', 'Computer Science'));
```

## Read Query

```
SELECT p.name FROM People p WHERE p IS OF (ONLY Student);
```

# Inheritance and Object Identity

- Object Identity (OID): Unique ID for objects.
- Each object has **a system-defined identifier** that remains constant, regardless of the object's values.
- Example:
  - An object of **Student** type inserted into a table will be assigned an internal OID.
  - You can use this OID **to track or reference the object** even if its attributes change.

# Inheritance and Object Identity

- Reference Types: Objects pointing to others using REF types in SQL.

```
CREATE TYPE Department AS (name TEXT);
```

```
CREATE TABLE Departments OF Department;
```

```
CREATE TYPE Instructor UNDER Person AS (dept REF(Department));
```

```
CREATE TABLE Instructors OF Instructor;
```

Querying with Dereferencing

```
SELECT i.name, DEREF(i.dept).name FROM Instructors i;
```

# Query Differences — OODB vs ORDB

## 1) Navigating Collections

### OODB: Implicit navigation

```
SELECT s.name  
FROM Course c, c.enrolledStudents s  
WHERE c.title = 'Database Systems';
```

- **enrolledStudents** is a list of object references, and the query language (e.g., OQL) handles iteration automatically.

### ORDB: Explicit unnesting required

```
SELECT s.name  
FROM Course c, UNNEST(c.enrolledStudents) AS s  
WHERE c.title = 'Database Systems';
```

- *You must manually flatten collection attributes like arrays or multisets.*

# Query Differences — OODB vs ORDB

## Dereferencing Object References

### OODB: Implicit dereferencing

```
SELECT s.name, s.advisor.name  
FROM Student s;
```

- The system understands `advisor` is a reference and automatically fetches attributes.

### ORDB: Explicit dereferencing using `DEREF()`

```
SELECT s.name, DEREF(s.advisor).name  
FROM Student s;
```

- Required to access referenced object's attributes.

# Query Differences — OODB vs ORDB

## Inheritance in Tables

**OODB:** Classes with inheritance — queries behave like object method calls

**ORDB:** Tables support type hierarchies with queries like:

SELECT \* FROM Person; -- includes Student and Instructor

SELECT \* FROM ONLY Person; -- excludes subtypes

SELECT p.name FROM People p WHERE p IS OF (Student); – polymorphism

# Query Differences — OODB vs ORDB

## Method Invocation

- **OODB:** Object methods are naturally available (e.g., `s.getAge()`)

**ORDB:** Only possible if the DBMS supports methods on UDTs

-- Hypothetical example

- `SELECT s.name, s.getAge() FROM Student s;`

# Pros and Cons of OODB and ORDB

## OODB

### Pros:

- **Seamless integration** with object-oriented programming.
- Natural modeling of complex and hierarchical data.
- Supports encapsulation and behavior through methods.
- Ideal for applications like CAD, simulations, and multimedia.

### Cons:

- **Limited standardization** and vendor support.
- **Less mature query optimization and indexing capabilities.**
- Steeper learning curve for developers not familiar with OOP.
- **Poor interoperability with existing relational tools.**

## ORDB

### Pros:

- **Combines strengths of relational and object-oriented models.**
- **SQL-based with extensions for structured and user-defined types.**
- Better support and adoption in commercial RDBMS (e.g., PostgreSQL, Oracle).
- Easier transition from legacy relational systems.

### Cons:

- More **complex SQL syntax and schema design.**
- **Method support varies across platforms and may be limited.**
- **Object behavior (methods) not as naturally integrated** as in OODBs.
- **Inheritance and polymorphism can be**

# Persistent Programming Languages

- A persistent programming language integrates data persistence directly into the language semantics.
- Objects or data structures created in the language can persist beyond program termination without requiring explicit serialization or database commands.

## key Features:

### Orthogonal Persistence

Any object, regardless of type, can be made persistent without changing how it is written.

### Transparency

No special code or syntax is required to store or retrieve persistent data—access is the same as for transient (in-memory) data.

### Type Safety

Type checking is maintained even for persistent data.

### Garbage Collection

Unused persistent objects may be collected over time.

- Advantages:
- 1) Simplifies code: less boilerplate for storing and retrieving data.
  - 2) Enhances productivity for complex and long-running applications (e.g., notebook)
  - 3) Reduces impedance mismatch between programming language and database, ala, direct object based communication rather than translation to SQL

# Temporal Databases

- Time-varying data support.
- Time dimensions:
  - Valid time: Real-world validity
  - Transaction time: Stored in DB
- SQL extensions: AS OF, VALID BETWEEN, TEMPORAL JOIN
- Applications: Auditing, compliance, time-series analysis.

```
SELECT salary  
FROM Employee  
WHERE empID = 101 AND PERIOD [2018-  
01-01 TO 2019-01-01] OVERLAPS  
employment_period;
```

## Real-World Examples:

- **Healthcare Records:** Track patient information over time, including diagnoses and treatments.
- **Financial Systems:** Monitor stock prices, account balances, and transactions over time.
- **Human Resources:** Maintain employee roles, salaries, and department changes historically.
- **Insurance Claims:** Track policy changes, claim timelines, and customer status.

# Spatial Databases

- Handle **geometric, topological, and geographic data.**
- **Use Cases:** GIS, mapping, location-based services, urban planning.
- **Features:**
  - Spatial data types (points, lines, polygons)
  - Spatial indexing (e.g., R-trees)
  - Spatial predicates (e.g., `ST_Within`, `ST_Distance`)
- **Example Query:**

```
SELECT name FROM Locations WHERE ST_Distance(geom,  
ST_Point(23.8, 90.4)) < 1000;
```

# Multimedia Databases

- Manage **image, audio, video, and associated metadata.**
- **Use Cases:** Digital libraries, surveillance, media asset management.
- **Features:**
  - Support for media types and attributes (resolution, codec, duration)
  - **Content-based retrieval** (e.g., image similarity)
  - **Storage and indexing of large binary objects** (BLOBs)
- **Example Query:**

```
SELECT title FROM Videos WHERE duration < 300 AND resolution =  
'1080p';
```

# Mobile Databases

- Designed for use on mobile or distributed devices.
- **Use Cases:** Offline applications, field data collection, mobile banking.
- **Features:**
  - Disconnected operation and synchronization
  - Location/context-aware data services
  - **Lightweight storage engines** (e.g., SQLite, Realm)
- **Example Scenario:**
  - A delivery app syncing route logs collected offline once network is available.

# Summary

- Modern DBs extend beyond traditional models.
- ORDB/OODB: Add structure and identity.
- Support for temporal, spatial, multimedia data.
- Persistence and mobility are shaping DBMS evolution.

# References

- Elmasri & Navathe – Fundamentals of Database Systems
- Silberschatz et al. – Database System Concepts
- Garcia-Molina et al. – Database Systems: The Complete Book
- C.J. Date et al. – Temporal Data & the Relational Model